

Index no:19APSE4288 Piraichselavan

```
C:\Users\pirai>docker pull bde2020/hadoop-namenode:latest
```

```
C:\Users\pirai>docker pull bde2020/hadoop-namenode:latest
latest: Pulling from bde2020/hadoop-namenode
Digest: sha256:fd74110805132d646cf6f12635efc0919e1fb2ac5bd376c536627fc261301e
Status: Image is up to date for bde2020/hadoop-namenode:latest
docker.io/bde2020/hadoop-namenode:latest

What's next:
  View a summary of image vulnerabilities and recommendations → docker scout quickview bde2020/hadoop-namenode:latest
```

C:\Users\pirai>docker images

C:\Users\pirai>docker images				
REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
dockerapp	1.1	27f84d9a8889	2 weeks ago	159MB
nginx	latest	3b25b682ea82	6 weeks ago	192MB
alpine	latest	91ef0af61f39	2 months ago	7.8MB
hello-world	latest	d2c94e258dcf	18 months ago	13.3kB
bde2020/hadoop-namenode	latest	b638307a2119	4 years ago	1.37GB

```
C:\Users\pirai>docker run bde2020/hadoop-namenode:latest
```

```
C:\Users\pirai>docker run bde2020/hadoop-namenode:latest
Configuring core
- Setting fs.defaultFS=hdfs://04b60d2522b1:8020
Configuring hdfs
- Setting dfs.namenode.name.dir=file:///hadoop/dfs/name
Configuring yarn
Configuring httpfs
Configuring kms
Configuring mapred
Configuring for multihommed network
Cluster name not specified
```

```
C:\Users\pirai>docker run -it --name hadoop-cluster202 -p 9872:9871 -p 8090:8089 -p 60070:50070  
bde2020/hadoop-namenode:latest /bin/bash
```

```
C:\Users\spiral\docker run -it --name hadoop-cluster202 -p 9872:9871 -p 8898:8899 -p 88970:59979 bde2020/hadoop-namenode:latest /bin/bash
Configuring core
- Setting fs.defaultFS=hdfs://80c224a83dd:8892
Configuring fs
  Setting dfs.namenode.name.dir=file:///hadoop/dfs/name
Configuring yarn
Configuring httpfs
Configuring kms
Configuring mapred
Configuring multihomed
Configuring for multihomed network
root@80c224a83dd:/# start-all.sh
Starting HDFS NameNode
[hadoop] Starting HDFS NameNode found
root@80c224a83dd:/# hdfs namenode -format
2024-11-18 08:41:58,228 INFO namenode.NameNode: STARTUP_MSG:
*****STARTUP_MSG: Starting NameNode*****
```

```
2024-11-18 08:41:59.347 INFO blockmanagement.BlockManagerSafeMode dfs.namenode.safemode.threshold-pct = 0.9999900128746033
2024-11-18 08:41:59.347 INFO blockmanagement.BlockManagerSafeMode dfs.namenode.safemode.min_datanodes = 0
2024-11-18 08:41:59.347 INFO blockmanagement.BlockManagerSafeMode dfs.namenode.safemode.extension = 30000
2024-11-18 08:41:59.349 INFO blockmanagement.BlockManager: defaultReplication = 3
2024-11-18 08:41:59.349 INFO blockmanagement.BlockManager: maxReplication = 512
2024-11-18 08:41:59.349 INFO blockmanagement.BlockManager: minReplication = 1
2024-11-18 08:41:59.349 INFO blockmanagement.BlockManager: maxReplicationStreams = 2
2024-11-18 08:41:59.350 INFO blockmanagement.BlockManager: redundancyRecheckInterval = 3000ms
2024-11-18 08:41:59.350 INFO blockmanagement.BlockManager: encryptDataTransfer = false
2024-11-18 08:41:59.351 INFO blockmanagement.BlockManager: maxNumBlocksToLog = 1000
2024-11-18 08:41:59.351 INFO namenode.FSDirectory: GLOBAL serial map: bits=29 maxEntries=536870911
2024-11-18 08:41:59.351 INFO namenode.FSDirectory: USER serial map: bits=29 maxEntries=16777215
2024-11-18 08:41:59.351 INFO namenode.FSDirectory: GROUP serial map: bits=24 maxEntries=16777215
2024-11-18 08:41:59.351 INFO namenode.FSDirectory: XATTR serial map: bits=24 maxEntries=16777215
2024-11-18 08:41:59.424 INFO util.GSet: Computing capacity for map INodeMap
2024-11-18 08:41:59.424 INFO util.GSet: VM type = 64-bit
2024-11-18 08:41:59.424 INFO util.GSet: 1.0% max memory 857 MB = 8.6 MB
2024-11-18 08:41:59.424 INFO util.GSet: capacity = 2^20 = 1048576 entries
2024-11-18 08:41:59.424 INFO namenode.FSDirectory: ACLs enabled: false
2024-11-18 08:41:59.424 INFO namenode.FSDirectory: XATTR ACL inheritance enabled? true
2024-11-18 08:41:59.426 INFO namenode.FSDirectory: XATTRs enabled: true
2024-11-18 08:41:59.427 INFO namenode.NameNode: Caching file names occurring more than 10 times
2024-11-18 08:41:59.440 INFO snapshot.SnapshotManager: Loaded config captureOpenFiles: false, skipCaptureAccessTimeOnlyChange: false, snapshotDiffAllowSnapRootDescendant: true, maxSnapshotLimit: 65536
2024-11-18 08:41:59.445 INFO snapshot.SnapshotManager: SkipList is disabled
2024-11-18 08:41:59.455 INFO util.GSet: Computing capacity for map cachedBlocks
2024-11-18 08:41:59.456 INFO util.GSet: VM type = 64-bit
2024-11-18 08:41:59.456 INFO util.GSet: 0.01% max memory 857 MB = 2.1 MB
2024-11-18 08:41:59.457 INFO util.GSet: capacity = 2^18 = 262144 entries
2024-11-18 08:41:59.477 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.window.num.buckets = 10
2024-11-18 08:41:59.478 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.num.users = 10
2024-11-18 08:41:59.478 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.windows.minutes = 1,5,25
2024-11-18 08:41:59.481 INFO namenode.FSNamesystem: Retry cache on namenode is enabled
2024-11-18 08:41:59.481 INFO namenode.FSNamesystem: Retry cache will use 0.03 of total heap and retry cache entry expiry time is 600000 millis
2024-11-18 08:41:59.481 INFO util.GSet: Computing capacity for map NameNodeRetryCache
2024-11-18 08:41:59.491 INFO util.GSet: VM type = 64-bit
2024-11-18 08:41:59.491 INFO util.GSet: 0.0001% max memory 857 MB = 64 bytes
2024-11-18 08:41:59.495 INFO util.GSet: capacity = 2^15 = 32768 entries
2024-11-18 08:41:59.601 INFO namenode.FSIImage: Allocated new BlockPoolId: BP-1418316472-172.17.0.2-1731919319553
2024-11-18 08:41:59.641 INFO common.Storage: Storage directory /hadoop/dfs/name has been successfully formatted.
2024-11-18 08:41:59.718 INFO namenode.FSIImageFormatProtobuf: Saving image file /hadoop/dfs/name/current/fsimage.checkpt_00000000000000000000000000000000 using no compression
2024-11-18 08:41:59.872 INFO namenode.FSIImageFormatProtobuf: Image file /hadoop/dfs/name/current/fsimage.checkpt_00000000000000000000000000000000 of size 399 bytes saved in 0 seconds .
2024-11-18 08:41:59.893 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0
2024-11-18 08:41:59.904 INFO namenode.FSIImage: FSIImageSaver clean checkpoint: txid=0 when meet shutdown.
2024-11-18 08:41:59.904 INFO namenode.NameNode: SHUTDOWN_MSG:
*****FSImageSaver clean checkpoint: txid=0 when meet shutdown.*****
SHUTDOWN_MSG: Shutting down NameNode at 80cd224a83d/172.17.0.2
*****
root@80cd224a83d:/# hdfs namenode &
[1] 203
```

```

hadoop-yarn-services-api-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-router-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-tests-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-common-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-services-core-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-applicationhistoryservice-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-timeline-pluginstorage-3.2.1.jar
STARTUP_MSG: build = https://gitbox.apache.org/repos/asf/hadoop.git -r b3ccb467e22ea829b3808f4b7b01d07e0bf3842; compiled by 'rohithsharmaks' on 2019-09-10T15:56Z
STARTUP_MSG: Java 1.8.0_232
*****
2024-11-18 08:43:37.206 INFO namenode.NameNode: registered UNIX signal handlers for [TERM, HUP, INT]
2024-11-18 08:43:37.366 INFO namenode.NameNode: createNameNode []
2024-11-18 08:43:37.587 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2024-11-18 08:43:37.752 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2024-11-18 08:43:37.752 INFO impl.MetricsSystemImpl: NameNode metrics system started
2024-11-18 08:43:37.811 INFO namenode.NameNodeUtils: fs.defaultFS is hdfs://80cd224a83dd:8020
2024-11-18 08:43:37.811 INFO namenode.NameNode: Clients should use 80cd224a83dd:8020 to access this namenode/service.
2024-11-18 08:43:38.001 INFO util.JvmPauseMonitor: Starting JVM pause monitor
2024-11-18 08:43:38.243 INFO util.WebResource: Started WebResource for dfs at: http://0.0.0.0:9870
2024-11-18 08:43:38.242 INFO util.LogFile: Logging initialized @255ms
2024-11-18 08:43:38.363 INFO server.AuthenticationFilter: Unable to initialize FileSignerSecretProvider, falling back to use random secrets.
2024-11-18 08:43:38.388 INFO http.HttpRequestLog: Http request log for http.requests.namenode is not defined
2024-11-18 08:43:38.403 INFO http.HttpServer2: Added global filter 'safety' (class=org.apache.hadoop.http.lib.StaticUserWebFilter$StaticUserFilter) to context hdfs
2024-11-18 08:43:38.403 INFO http.HttpServer2: Added filter static_user_filter (class=org.apache.hadoop.http.lib.StaticUserWebFilter$StaticUserFilter) to context logs
2024-11-18 08:43:38.404 INFO http.HttpServer2: Added filter static_user_filter (class=org.apache.hadoop.http.lib.StaticUserWebFilter$StaticUserFilter) to context static
2024-11-18 08:43:38.404 INFO http.HttpServer2: Added filter 'org.apache.hadoop.hdfs.web.AuthFilter' (class=org.apache.hadoop.hdfs.web.AuthFilter)
2024-11-18 08:43:38.404 INFO http.HttpServer2: addAddressResourcePackage: package name=org.apache.hadoop.hdfs.server.namenode.web.Resources; org.apache.hadoop.hdfs.web.Resources, pathSpec=/webhdfs/*
2024-11-18 08:43:38.473 INFO http.HttpServer2: Jetty bound to port 9870
2024-11-18 08:43:38.475 INFO server.Server: jetty-9.3.24.v20180605, build timestamp: 2018-06-05T17:11:56Z, git hash: 84285aa28f11a4f31f2a3b8d1bba2cc8ab69827
2024-11-18 08:43:38.526 INFO handler.ContextHandler: Started o.e.j.s.ServletContextHandler@7db07e3c{/logs,file:///opt/hadoop-3.2.1/logs,AVAILABLE}
2024-11-18 08:43:38.527 INFO handler.ContextHandler: Started o.e.j.s.ServletContextHandler@4b741d6d{/static,file:///opt/hadoop-3.2.1/share/hadoop/hdfs/webapps/static/,AVAILABLE}
2024-11-18 08:43:38.681 INFO handler.ContextHandler: Started o.e.j.w.WebAppContext@7486b455{/file:///opt/hadoop-3.2.1/share/hadoop/hdfs/webapps/hdfs/,AVAILABLE}{/hdfs}
2024-11-18 08:43:38.713 INFO server.AbstractConnector: Started ServerConnector@700fb071{HTTP/1.1}{0.0.0.0:9870}
2024-11-18 08:43:38.713 INFO server.Server: Started @3927ms
2024-11-18 08:43:38.713 INFO server.Server: Started @3927ms
2024-11-18 08:43:38.969 WARN namenode.FSNamesystem: Only one image storage directory (dfs.namenode.name.dir) configured. Beware of data loss due to lack of redundant storage directories!
2024-11-18 08:43:38.969 WARN namenode.FSNamesystem: Only one namespace edits storage directory (dfs.namenode.edits.dir) configured. Beware of data loss due to lack of redundant storage directories!
2024-11-18 08:43:39.028 INFO namenode.FSEditLog: Edit logging is async:true
2024-11-18 08:43:39.083 INFO namenode.FSNamesystem: KeyProvider: null
2024-11-18 08:43:39.083 INFO namenode.FSNamesystem: fsLock is fair: true
2024-11-18 08:43:39.085 INFO namenode.FSNamesystem: Detailed lock hold time metrics enabled: false
2024-11-18 08:43:39.085 INFO namenode.FSNamesystem: fsOwner = root (auth:SIMPLE)
2024-11-18 08:43:39.089 INFO namenode.FSNamesystem: supergroup = supergroup
2024-11-18 08:43:39.099 INFO namenode.FSNamesystem: isIPPermissionEnabled = true
2024-11-18 08:43:39.099 INFO namenode.FSNamesystem: HA Enabled: false
2024-11-18 08:43:39.146 INFO common.Util: dfs.datanode.fileio.profiling.sampling.percentage set to 0. Disabling file IO profiling
2024-11-18 08:43:39.157 INFO blockmanagement.DatanodeManager: dfs.block.invalidate.limit: configured=1080, counted=60, effected=1080
2024-11-18 08:43:39.157 INFO blockmanagement.DatanodeManager: dfs.namenode.datanode.registration.ip-hostname-check=true
2024-11-18 08:43:39.161 INFO blockmanagement.BlockManager: dfs.namenode.startup.delay.block.deletion.sec is set to 000:00:00.00.00.00
2024-11-18 08:43:39.161 INFO blockmanagement.BlockManager: The block deletion will start around 2024 Nov 18 08:43:39
```

```

2024-11-18 08:43:39.371 INFO namenode.FSNamesystem: Retry cache will use 0.03 of total heap and retry cache entry expiry time is 600000 millis
2024-11-18 08:43:39.381 INFO util.GSet: Computing capacity for map NameNodeRetryCache
2024-11-18 08:43:39.381 INFO util.GSet: VM type = 64-bit
2024-11-18 08:43:39.383 INFO util.GSet: max memory 857 MB = 263.3 KB
2024-11-18 08:43:39.383 INFO util.GSet: capacity = 2^15 = 32768 entries
2024-11-18 08:43:39.445 INFO common.Storage: Lock on /hadoop/dfs/name/inuse.lock acquired by nodename 203080cd224a83dd
2024-11-18 08:43:39.497 INFO namenode.FileJournalManager: Recovering unfinalized segments in /hadoop/dfs/name/current
2024-11-18 08:43:39.497 INFO namenode.FSImage: No edit log streams selected.
2024-11-18 08:43:39.498 INFO namenode.FSImage: Planning to load image from /hadoop/dfs/name/current/fsimage_00000000000000000000, cpktTxId=00000000000000000000
2024-11-18 08:43:39.593 INFO namenode.FSImageFormatPBINode: Loading 1 INodes.
2024-11-18 08:43:39.675 INFO namenode.FSImageFormatProtobuf: Loaded FSImage in 0 seconds.
2024-11-18 08:43:39.675 INFO namenode.FSImageFormatProtobuf: Need to save fs image? false (staleImage=false, haEnabled=false, isRollingUpgrade=false)
2024-11-18 08:43:39.682 INFO namenode.FSEditLog: Starting log segment at 1
2024-11-18 08:43:39.917 INFO namenode.NameCache: initialized with 0 entries 0 lookups
2024-11-18 08:43:39.917 INFO namenode.FSNamesystem: Finished loading FSImage in 513 msecs
2024-11-18 08:43:40.266 INFO namenode.NameNode: RPC server is binding to 0.0.0.0:8020
2024-11-18 08:43:40.298 INFO ipc.CallQueueManager: Using callQueue: class java.util.concurrent.LinkedBlockingQueue, queueCapacity: 1000, scheduler: class org.apache.hadoop.ipc.IpcScheduler ipcBackoff: false.
2024-11-18 08:43:40.312 INFO ipc.Server: Starting Socket Reader #1 for port 8020
2024-11-18 08:43:40.571 INFO namenode.FSNamesystemState: Registered FSNamesystemState, ReplicatedBlocksState and ECBBlockGroupsState MBeans.
2024-11-18 08:43:40.590 INFO namenode.LeaseManager: Number of blocks under construction: 0
2024-11-18 08:43:40.601 INFO blockmanagement.BlockManager: initializing replication queues
2024-11-18 08:43:40.602 INFO namenode.StateChange: STATE* Leaving safe mode after 0 secs
2024-11-18 08:43:40.602 INFO namenode.StateChange: STATE* Network topology has 0 racks and 0 datanodes
2024-11-18 08:43:40.602 INFO dfs.StateChange: STATE* UnderReplicatedBlocks has 0 blocks
2024-11-18 08:43:40.613 INFO blockmanagement.BlockManager: Total number of blocks = 0
2024-11-18 08:43:40.613 INFO blockmanagement.BlockManager: Number of invalid blocks = 0
2024-11-18 08:43:40.613 INFO blockmanagement.BlockManager: Number of under-replicated blocks = 0
2024-11-18 08:43:40.613 INFO blockmanagement.BlockManager: Number of over-replicated blocks = 0
2024-11-18 08:43:40.613 INFO blockmanagement.BlockManager: Number of blocks being written = 0
2024-11-18 08:43:40.613 INFO hdfs.StateChange: STATE* Replication Queue initialization scan for invalid, over- and under-replicated blocks completed in 11 msec
2024-11-18 08:43:40.641 INFO ipc.Server: IPC Server Responder: starting
2024-11-18 08:43:40.645 INFO ipc.Server: IPC Server Listener on 8020: starting
2024-11-18 08:43:40.659 INFO namenode.NameNode: NameNode RPC up at: 80cd224a83dd/172.17.0.2:8020
2024-11-18 08:43:40.663 INFO namenode.FSNamesystem: Starting services required for active state
2024-11-18 08:43:40.664 INFO namenode.FSDirectory: Initializing quota with 4 thread(s)
2024-11-18 08:43:40.676 INFO namenode.FSDirectory: Quota initialization completed in 12 milliseconds
name_space=1
storage_space=0
storage_types=RAM,DISK=0, SSD=0, DISK=0, ARCHIVE=0, PROVIDED=0
2024-11-18 08:43:40.687 INFO blockmanagement.CacheReplicationMonitor: Starting CacheReplicationMonitor with interval 30000 milliseconds
```

```
root@80cd224a83dd:/# hdfs datanode &
```

```
2024-11-18 08:45:53.358 INFO common.Storage: Storage directory with location [/DISK/file:/tmp/hadoop-root/dfs/data] is not formatted for namespace 671648660. Formatting...
2024-11-18 08:45:53.361 INFO common.Storage: Generated new StorageID DS-eb48ec27-c2d6-4096-b47a-586161493f26 for directory /tmp/hadoop-root/dfs/data
2024-11-18 08:45:53.391 INFO common.Storage: Analyzing storage directores for bpid BP-1418316472-172.17.0.2-1731919319553
2024-11-18 08:45:53.391 INFO common.Storage: Locking is disabled for /tmp/hadoop-root/dfs/data/current/BP-1418316472-172.17.0.2-1731919319553
2024-11-18 08:45:53.392 INFO common.Storage: Block pool storage directory for location [/DISK/file:/tmp/hadoop-root/dfs/data and block pool id BP-1418316472-172.17.0.2-1731919319553 is not formatted. Formatting ...
2024-11-18 08:45:53.392 INFO common.Storage: Formating block pool BP-1418316472-172.17.0.2-1731919319553 directory /tmp/hadoop-root/dfs/data/current/BP-1418316472-172.17.0.2-1731919319553
2024-11-18 08:45:53.399 INFO datanode.DataNode: Setting up storage: nsid=671648660;bpid=BP-1418316472-172.17.0.2-1731919319553;lnsInfo=fhv=65;cid=CID-e694942a-0673-423f-824a-a2a0386365ec;nsid=671648660;bpid=BP-1418316472-172.17.0.2-1731919319553;dnurl=null
2024-11-18 08:45:53.401 INFO datanode.DataNode: Generated and persisted new Datanode UUID 0873a3b0-25df-423e-9395-73df960fd4d5
2024-11-18 08:45:53.547 INFO impl.FsDatasetImpl: Added new volume: DS-eb48ec27-c2d6-4096-b47a-586161493f26
2024-11-18 08:45:53.547 INFO impl.FsDatasetImpl: Added volume - [/DISK/file:/tmp/hadoop-root/dfs/data, StorageType: DISK
2024-11-18 08:45:53.551 INFO impl.FsDatasetImpl: Registered FSDataState Mbean
2024-11-18 08:45:53.551 INFO checked.ThrottledAsyncChecker: Scheduling a check for /tmp/hadoop-root/dfs/data
2024-11-18 08:45:53.556 INFO impl.FsDatasetImpl: Starting a check for volume /tmp/hadoop-root/dfs/data
2024-11-18 08:45:53.568 INFO impl.FsDatasetImpl: Adding block pool BP-1418316472-172.17.0.2-1731919319553
2024-11-18 08:45:53.570 INFO impl.FsDatasetImpl: Scanning block pool BP-1418316472-172.17.0.2-1731919319553 on volume /tmp/hadoop-root/dfs/data...
2024-11-18 08:45:53.610 INFO impl.FsDatasetImpl: Total time to scan block pool BP-1418316472-172.17.0.2-1731919319553 on /tmp/hadoop-root/dfs/data: 46ms
2024-11-18 08:45:53.616 INFO impl.FsDatasetImpl: Total time to scan all replicas for block pool BP-1418316472-172.17.0.2-1731919319553: 48ms
2024-11-18 08:45:53.619 INFO impl.FsDatasetImpl: Adding replicas to map for block pool BP-1418316472-172.17.0.2-1731919319553 on volume /tmp/hadoop-root/dfs/data...
2024-11-18 08:45:53.619 INFO impl.BlockPoolSlice.Replica.Cache: file: /tmp/hadoop-root/dfs/data/current/BP-1418316472-172.17.0.2-1731919319553/current/replicas doesn't exist
2024-11-18 08:45:53.620 INFO impl.FsDatasetImpl: Total time to add replicas to map for block pool BP-1418316472-172.17.0.2-1731919319553 on volume /tmp/hadoop-root/dfs/data: 4ms
2024-11-18 08:45:53.620 INFO impl.FsDatasetImpl: Total time to add all replicas to map for block pool BP-1418316472-172.17.0.2-1731919319553: 4ms
2024-11-18 08:45:53.628 INFO datanode.VolumeScanner: Now scanning bpid BP-1418316472-172.17.0.2-1731919319553 on volume /tmp/hadoop-root/dfs/data
2024-11-18 08:45:53.628 INFO datanode.VolumeScanner: VolumeScanner(/tmp/hadoop-root/dfs/data, DS-eb48ec27-c2d6-4096-b47a-586161493f26): finished scanning block pool BP-1418316472-172.17.0.2-1731919319553
2024-11-18 08:45:53.644 INFO datanode.VolumeScanner: VolumeScanner(/tmp/hadoop-root/dfs/data, DS-eb48ec27-c2d6-4096-b47a-586161493f26): no suitable block pools found to scan. Waiting 181439981 ms.
2024-11-18 08:45:53.646 INFO datanode.DirectoryScanner: Periodic Directory Tree Visibility Scanner scan starting at 11/18/24 1:50 PM with interval of 21600000ms
2024-11-18 08:45:53.656 INFO datanode.DataNode: Block pool BP-1418316472-172.17.0.2-1731919319553 (Datanode Uuid 0873a3b0-25df-423e-9395-73df960fd4d5) service to 80cd224a83dd/172.17.0.2:80
2024-11-18 08:45:53.656 INFO datanode.DataNode: Beginning handshake with master
2024-11-18 08:45:53.713 INFO common.Storage: StateChange: BLOCS< registerDatanode: from DatanodeScanner [172.17.0.2-9866, datanodeUuid=0873a3b0-25df-423e-9395-73df960fd4d5, infoPort=9864, infoSecurePort=9867, storageInfo=fhv=57;cid:CID-e694942a-0673-423f-824a-d3a036365ec;nsid=671648660;c=1731919319553] storage 0873a3b0-25df-423e-9395-73df960fd4d5
2024-11-18 08:45:53.715 INFO net.NetworkTopology: Adding a new node: /default-rack[/172.17.0.2-9866]
2024-11-18 08:45:53.715 INFO blockmanagement.BlockReportLeaseManager: Registered DN 0873a3b0-25df-423e-9395-73df960fd4d5 (17.17.0.2-9866).
2024-11-18 08:45:53.732 INFO datanode.DataNode: Block pool Block pool BP-1418316472-172.17.0.2-1731919319553 (Datanode Uuid 0873a3b0-25df-423e-9395-73df960fd4d5) service to 80cd224a83dd/172.17.0.2:8020 successfully registered with NN
2024-11-18 08:45:53.732 INFO datanode.DataNode: For namenode 80cd224a83dd/172.17.0.2:8020 using BLOCREPORT_INTERVAL of 1000msec Initial delay: 0sec; heartbeatInterval=8000ms
2024-11-18 08:45:53.800 INFO blockmanagement.DatanodeDescriptor: Adding new storage ID DS-eb48ec27-c2d6-4096-b47a-586161493f26 for DN 172.17.0.2-9866
2024-11-18 08:45:53.800 INFO blockmanagement.DatanodeDescriptor: Adding new storage ID DS-eb48ec27-c2d6-4096-b47a-586161493f26 for DN 172.17.0.2-9866
2024-11-18 08:45:53.800 INFO BlockStateChange: BLOCS< processReport 0x52c7ba89e78c602: from storage DS-eb48ec27-c2d6-4096-b47a-586161493f26 node DatanodeScanner [172.17.0.2-9866, dataNodeUuid=0873a3b0-25df-423e-9395-73df960fd4d5, infoPort=9864, infoSecurePort=0, ipPort=9867, storageInfo=fhv=57;cid:CID-e694942a-0673-423f-824a-d3a036365ec;nsid=671648660;c=1731919319553], blocks: 0, hasStaleStorage: false, processing time: 3 msecs, invalidatedBlocks: 0
2024-11-18 08:45:53.917 INFO datanode.DataNode: Successfully sent block report 0x52c7ba89e78c602, containing 1 storage report(s), of which we sent 1. The reports had 0 total blocks and u
sed 1 RPC(s). This took 3 msec to generate and 56 msecs for RPC and NN processing. Got back one command: FinalizeCommand/5.
2024-11-18 08:45:53.918 INFO datanode.DataNode: Got finalize command for block pool BP-1418316472-172.17.0.2-1731919319553
```

```
root@80cd224a83dd:/# yarn nodemanager &
```

```
2024-11-18 08:47:13,749 WARN monitor.ContainerMonitorImpl: NodeManager configured with 8 G physical memory allocated to containers, which is more than 80% of the total physical memory available (3.8 G). Thrashing might happen.
2024-11-18 08:47:13,750 INFO containermanager.ContainerManagerImpl: Not a recoverable state store. Nothing to recover.
2024-11-18 08:47:13,781 INFO conf.Configuration: resource-types.xml not found
2024-11-18 08:47:13,794 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2024-11-18 08:47:13,801 INFO conf.Configuration: node-resources.xml not found
2024-11-18 08:47:13,802 INFO conf.Configuration: resources.xml not found
2024-11-18 08:47:13,810 INFO nodemanager.NodeStatusUpdaterImpl: NodeManager resources is set to: <memory:8192, vCores:8>
2024-11-18 08:47:13,824 INFO nodemanager.NodeStatusUpdaterImpl: Initialized nodemanager with : physical-memory=8192 virtual-memory=17284 virtual-cores=8
2024-11-18 08:47:13,944 INFO ipc.CallQueueManager: Using callQueue: class java.util.concurrent.LinkedBlockingQueue, queueCapacity: 2080, scheduler: class org.apache.hadoop.ipc.DefaultRpcScheduler, ipcBackoff: false.
2024-11-18 08:47:14,002 INFO ipc.Server: Starting Socket Reader #1 for port 0
2024-11-18 08:47:14,482 INFO pb.RpcServerFactoryPBImpl: Adding protocol org.apache.hadoop.yarn.api.ContainerManagementProtocolPB to the server
2024-11-18 08:47:14,483 INFO ipc.Server: IPC Server Responder: starting
2024-11-18 08:47:14,502 INFO ipc.Server: IPC Server listener on 0: starting
2024-11-18 08:47:14,502 INFO security.NMContainerTokenSecretManager: Updating node address : 80cd224a83d3d:44213
2024-11-18 08:47:14,514 INFO ipc.CallQueueManager: Using callQueue: class java.util.concurrent.LinkedBlockingQueue, queueCapacity: 500, scheduler: class org.apache.hadoop.ipc.DefaultRpcScheduler, ipcBackoff: false.
2024-11-18 08:47:14,516 INFO ipc.Server: Starting Socket Reader #1 for port 8940
2024-11-18 08:47:14,521 INFO pb.RpcServerFactoryPBImpl: Adding protocol org.apache.hadoop.yarn.server.nodemanager.api.LocalizationProtocolPB to the server
2024-11-18 08:47:14,523 INFO ipc.Server: IPC Server Responder: starting
2024-11-18 08:47:14,523 INFO ipc.Server: IPC Server listener on 8940: starting
2024-11-18 08:47:14,526 INFO localization.LocalizationManagerImpl: LocalizationManager started on port 8940
2024-11-18 08:47:14,530 INFO containermanager.ContainerManagerImpl: ContainerManager started at 80cd224a83d3d/172.17.0.2:44213
2024-11-18 08:47:14,530 INFO containermanager.ContainerManagerImpl: ContainerManager bound to 0.0.0.0/0.0.0.0:8
2024-11-18 08:47:14,534 INFO NMLogAggregationStatusTracker: Log Aggregation is disabled.So is the LogAggregationStatusTracker.
2024-11-18 08:47:14,536 INFO webapp.WebServer: Instantiating NWWebApp at 0.0.0.0:8942
2024-11-18 08:47:14,627 INFO util.Log: Logging initialized @3330ms.
2024-11-18 08:47:14,828 INFO server.AuthenticationFilter: Unable to initialize FileSignerSecretProvider, falling back to use random secrets.
2024-11-18 08:47:14,832 INFO http.HttpRequestLog: HTTP request log for http://requests.nodemanager is not defined
2024-11-18 08:47:14,840 INFO http.HttpServer: Global filter 'staticUserFilter' (class=org.apache.hadoop.http.lib.StaticUserFilter)
2024-11-18 08:47:14,841 INFO http.HttpServer: Added filter static_user_filter (class=org.apache.hadoop.http.lib.StaticUserWebFilter$StaticUserFilter) to context node
2024-11-18 08:47:14,842 INFO http.HttpServer: Added filter static_user_filter (class=org.apache.hadoop.http.lib.StaticUserWebFilter$StaticUserFilter) to context logs
2024-11-18 08:47:14,845 INFO http.HttpServer: Added filter static_user_filter (class=org.apache.hadoop.http.lib.StaticUserWebFilter$StaticUserFilter) to context static
2024-11-18 08:47:14,846 INFO http.HttpServer: Added filter authentication (class=org.apache.hadoop.security.authentication.server.AuthenticationFilter) to context node
2024-11-18 08:47:14,847 INFO http.HttpServer: Added filter authentication (class=org.apache.hadoop.security.authentication.server.AuthenticationFilter) to context logs
2024-11-18 08:47:14,848 INFO http.HttpServer: Added filter authentication (class=org.apache.hadoop.security.authentication.server.AuthenticationFilter) to context static
2024-11-18 08:47:14,850 INFO http.HttpServer: Adding path spec: /node/*
2024-11-18 08:47:14,851 INFO http.HttpServer: Adding path spec: /status/*
2024-11-18 08:47:14,852 INFO http.HttpServer: Webapp modules loaded: guice modules
2024-11-18 08:47:15,368 INFO http.HttpServer: Jetty bound to port 8942
2024-11-18 08:47:15,376 INFO server.Server: jetty-9.3.24.v20180605, build timestamp: 2018-06-05T11:56Z, git hash: 84205aa28f11a4f31f2a3b86d1bba2ccBab69827
2024-11-18 08:47:15,420 INFO server.AuthenticationFilter: Unable to initialize FileSignerSecretProvider, falling back to use random secrets.
2024-11-18 08:47:15,424 INFO handler.ContextHandler: Started o.e.j.s.ServletContextHandler@5a9d6f02/{logs,file:///opt/hadoop-3.2.1/logs/,AVAILABLE}
2024-11-18 08:47:15,431 INFO handler.ContextHandler: Started o.e.j.s.ServletContextHandler@6097f88e/{static,jar,file:///opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-common-3.2.1.jar!/webapps/static/,AVAILABLE}
2024-11-18 08:47:15,432 INFO am.sun.jersey.guice.spi.container.GuiceComponentProviderFactory: Registering org.apache.hadoop.yarn.server.nodemanager.webapp.NWWebServices as a root resource class
INFO: Registering org.apache.hadoop.yarn.webapp.GuiceExceptionHandler as provider class
Nov 18 2024 8:47:15 AM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory register
INFO: Registering org.apache.hadoop.yarn.server.nodemanager.webapp.NWWebServices as a root resource class
Nov 18 2024 8:47:15 AM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory register
```

```
2024-11-18 08:47:15.362 INFO webapp.WebApps: Registered mewapp core modules
2024-11-18 08:47:15.368 INFO webapp.HttpServer2: Jetty bound to port 8084
2024-11-18 08:47:15.370 INFO server.Server: Jetty-9.3.24.v20180605, build timestamp: 2018-06-05T17:11:56Z, git hash: 84205aa28f1la4f31f2a3b86d1bba2cc8ab69827
2024-11-18 08:47:15.426 INFO server.AuthenticationFilter: Unable to initialize FileSignerSecretProvider, falling back to use random secrets.
2024-11-18 08:47:15.428 INFO handler.ContextHandler: Started o.e.j.s.ServletContextHandler@6d9ff6f02/{logs,file:///opt/hadoop-3.2.1/logs/,AVAILABLE}
2024-11-18 08:47:15.431 INFO handler.ContextHandler: Started o.e.j.s.ServletContextHandler@6d9ff7a80/{static,jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-common-3.2.1.jar!/webapps/static/AVAILABLE}
Nov 18, 2024 8:47:15 AM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory register
INFO: Registering org.apache.hadoop.yarn.server.nodemanager.webapp.WNWebServices as a root resource class
Nov 18, 2024 8:47:15 AM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory register
INFO: Registering org.apache.hadoop.yarn.webapp.GenericExceptionHandler as a provider class
Nov 18, 2024 8:47:15 AM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory register
INFO: Registering org.apache.hadoop.yarn.server.nodemanager.webapp.JAXBContextResolver as a provider class
Nov 18, 2024 8:47:15 AM com.sun.jersey.server.impl.application.WebApplicationImpl _initiate
INFO: Initiating Jersey application, version 'Jersey' 1.19 02/11/2015 03:25 AM'
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.JAXBContextResolver to GuiceManagedComponentProvider with the scope "Singleton"
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.WNWebServices to GuiceManagedComponentProvider with the scope "Singleton"
INFO: Binding org.apache.hadoop.yarn.webapp.GenericExceptionHandler to GuiceManagedComponentProvider with the scope "Singleton"
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.WNWebServices to GuiceManagedComponentProvider with the scope "Singleton"
Nov 18, 2024 8:47:16 AM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.JAXBContextResolver to GuiceManagedComponentProvider with the scope "Singleton"
Nov 18, 2024 8:47:16 AM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.WNWebServices to GuiceManagedComponentProvider with the scope "Singleton"
Nov 18, 2024 8:47:16.562 INFO handler.ContextHandler: Started o.e.j.s.WebAppContext@3533d1f61{/file:///tmp/jetty-0.0.0.0-8042-node_--any-_7419608939530946615.dir/webapp/,AVAILABLE}{/node}
2024-11-18 08:47:16.575 INFO server.AbstractConnector: Started ServerConnector@84d96250[HTTP/1.1,[http://1.1]{0.0.0.0:8042}]
2024-11-18 08:47:16.575 INFO server.Server: Started 65279ms
2024-11-18 08:47:16.576 INFO webapp.WebApps: Web app node started at 8042
2024-11-18 08:47:16.579 INFO nodemanager.NodeStatusUpdaterImpl: Node ID assigned is : 80cd224a83dd:44213
2024-11-18 08:47:16.580 INFO util.JVMPauseMonitor: Starting JVM pause monitor
2024-11-18 08:47:16.600 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8081
2024-11-18 08:47:16.657 INFO nodemanager.NodeStatusUpdaterImpl: Ending port 0 RM container statuses: []
2024-11-18 08:47:16.699 INFO nodemanager.NodeStatusUpdaterImpl: Registering with RM using containers: []
2024-11-18 08:47:17.822 INFO ipc.Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8081. Already tried 0 time(s); retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISSECONDS)
2024-11-18 08:47:18.824 INFO ipc.Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8081. Already tried 1 time(s); retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISSECONDS)
2024-11-18 08:47:19.826 INFO ipc.Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8081. Already tried 2 time(s); retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISSECONDS)
2024-11-18 08:47:20.827 INFO ipc.Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8081. Already tried 3 time(s); retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISSECONDS)
2024-11-18 08:47:21.829 INFO ipc.Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8081. Already tried 4 time(s); retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISSECONDS)
2024-11-18 08:47:22.830 INFO ipc.Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8081. Already tried 5 time(s); retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISSECONDS)
2024-11-18 08:47:23.831 INFO ipc.Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8081. Already tried 6 time(s); retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISSECONDS)
2024-11-18 08:47:24.833 INFO ipc.Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8081. Already tried 7 time(s); retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISSECONDS)
2024-11-18 08:47:25.839 INFO ipc.Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8081. Already tried 8 time(s); retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISSECONDS)
2024-11-18 08:47:26.843 INFO ipc.Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8081. Already tried 9 time(s); retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISSECONDS)
```

```
root@80cd224a83dd:/# yarn resourcemanager &
```

```
root@80cd224a83dd:/# hdfs dfs -mkdir -p /user/hadoop/input
```

```
root@80cd224a83dd:/# hdfs dfs -put $HADOOP_HOME/etc/hadoop/*.xml /user/hadoop/input
```

```
root@00cd224a83dd:~/hdfs dfs -mkdir -p /user/hadoop/input
2024-11-18 08:48:32,841 INFO namenode.FSEditLog: Number of transactions: 4 Total time for transactions(ms): 30 Number of transactions batched in Syncs: 0 Number of syncs: 2 syncTimes(ms): 23
root@00cd224a83dd:~/hdfs dfs -put $HADOOP_HOME/etc/hadoop/*_xml ./user/hadoop/input
2024-11-18 08:49:03 437 INFO hdfs.StateChange: BLOCK* allocate blk_1073741825_1001, replicas=>172.17.0.2:9866 for /user/hadoop/input/capacity-scheduler.xml._COPYING_
2024-11-18 08:49:09 532 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhost:trusted = false, remoteHostTrusted = false
2024-11-18 08:49:09 699 INFO dataNode.DataNode: Receiving BP->1418316472->172.17.0.2:1731919319553:blk_1073741825_1001 src: /172.17.0.2:41292 dest: /172.17.0.2:9866
2024-11-18 08:49:09 837 INFO dataNode.DataNode.clienttrace: src: /172.17.0.2:41282 dest: /172.17.0.2:9866 bytes: 8266 op: HDFS_WRITE clid: DFSClient_NONMAPREDUCE_-1824364460_1, offset: 0, srvID: 0, duration(ms): 429, blockID: 1073741825_1001, type:LAST_IN_PIPELINE
2024-11-18 08:49:09 837 INFO dataNode.DataNode: PacketResponder: BP->1418316472->172.17.0.2:1731919319553:blk_1073741825_1001, type:LAST_IN_PIPELINE terminating
2024-11-18 08:49:09 863 INFO namenode.FSNamesystem: BLOCK< blk_1073741825_1001 is COMMITTED but not COMPLETE(nuNodes= 0 < minimum = 1) in file /user/hadoop/input/capacity-scheduler.xml._COPYING_
2024-11-18 08:49:09 863 INFO hdfs.StateChange: DIR* completeFile: /user/hadoop/input/capacity-scheduler.xml._COPYING_ is closed by DFSClient_NONMAPREDUCE_-1824364460_1
2024-11-18 08:49:09 331 INFO hdfs.StateChange: BLOCK* allocate blk_1073741826_1002, replicas=>172.17.0.2:9866 for /user/hadoop/input/core-site.xml._COPYING_
2024-11-18 08:49:09 338 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhost:trusted = false, remoteHostTrusted = false
2024-11-18 08:49:09 344 INFO dataNode.DataNode: Receiving BP->1418316472->172.17.0.2:1731919319553:blk_1073741826_1002 src: /172.17.0.2:41212 dest: /172.17.0.2:9866
2024-11-18 08:49:09 350 INFO dataNode.DataNode.clienttrace: src: /172.17.0.2:41212 dest: /172.17.0.2:9866 bytes: 860, op: HDFS_WRITE, clid: DFSClient_NONMAPREDUCE_-1824364460_1, offset: 0, srvID: 0, duration(ms): 426, blockID: 1073741826_1002, type:LAST_IN_PIPELINE
2024-11-18 08:49:09 355 INFO hdfs.StateChange: DIR* completeFile: /user/hadoop/input/core-site.xml._COPYING_ is closed by DFSClient_NONMAPREDUCE_-1824364460_1
2024-11-18 08:49:09 385 INFO hdfs.StateChange: BLOCK* allocate blk_1073741827_1003, replicas=>172.17.0.2:9866 for /user/hadoop/input/hadoop-policy.xml._COPYING_
2024-11-18 08:49:09 391 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhost:trusted = false, remoteHostTrusted = false
2024-11-18 08:49:09 393 INFO dataNode.DataNode: Receiving BP->1418316472->172.17.0.2:1731919319553:blk_1073741827_1003, duration(ms): 1003, blockID: 1073741827_1003, type:LAST_IN_PIPELINE terminating
2024-11-18 08:49:09 404 INFO dataNode.DataNode.clienttrace: src: /172.17.0.2:41214 dest: /172.17.0.2:9866 bytes: 11392, op: HDFS_WRITE, clid: DFSClient_NONMAPREDUCE_-1824364460_1, offset: 0, srvID: 0, duration(ms): 350, blockID: 1073741827_1003, type:LAST_IN_PIPELINE
2024-11-18 08:49:09 404 INFO dataNode.DataNode: PacketResponder: BP->1418316472->172.17.0.2:1731919319553:blk_1073741827_1003, duration(ms): 6974505
2024-11-18 08:49:09 437 INFO hdfs.StateChange: DIR* completeFile: /user/hadoop/input/hadoop-policy.xml._COPYING_ is closed by DFSClient_NONMAPREDUCE_-1824364460_1
2024-11-18 08:49:09 437 INFO hdfs.StateChange: BLOCK* allocate blk_1073741828_1004, replicas=>172.17.0.2:9866 for /user/hadoop/input/hdfs-site.xml._COPYING_
2024-11-18 08:49:09 443 INFO dataNode.DataNode: Receiving BP->1418316472->172.17.0.2:1731919319553:blk_1073741828_1004, duration(ms): 1004, blockID: 1073741828_1004, type:LAST_IN_PIPELINE terminating
2024-11-18 08:49:09 443 INFO dataNode.DataNode.clienttrace: src: /172.17.0.2:41226 dest: /172.17.0.2:9866 bytes: 1385, op: HDFS_WRITE, clid: DFSClient_NONMAPREDUCE_-1824364460_1, offset: 0, duration(ms): 5625391
2024-11-18 08:49:09 456 INFO dataNode.DataNode: PacketResponder: BP->1418316472->172.17.0.2:1731919319553:blk_1073741828_1004, duration(ms): 5625391
2024-11-18 08:49:09 456 INFO dataNode.DataNode: Receiving BP->1418316472->172.17.0.2:1731919319553:blk_1073741828_1004, duration(ms): 1004, blockID: 1073741828_1004, type:LAST_IN_PIPELINE terminating
2024-11-18 08:49:09 459 INFO hdfs.StateChange: DIR* completeFile: /user/hadoop/input/hdfs-site.xml._COPYING_ is closed by DFSClient_NONMAPREDUCE_-1824364460_1
2024-11-18 08:49:09 488 INFO hdfs.StateChange: BLOCK* allocate blk_1073741829_1005, replicas=>172.17.0.2:9866 for /user/hadoop/input/httpfs-site.xml._COPYING_
2024-11-18 08:49:09 491 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhost:trusted = false, remoteHostTrusted = false
2024-11-18 08:49:09 497 INFO dataNode.DataNode: Receiving BP->1418316472->172.17.0.2:1731919319553:blk_1073741829_1005 src: /172.17.0.2:41238 dest: /172.17.0.2:9866
2024-11-18 08:49:09 504 INFO dataNode.DataNode.clienttrace: src: /172.17.0.2:41238 dest: /172.17.0.2:9866 bytes: 620, op: HDFS_WRITE, clid: DFSClient_NONMAPREDUCE_-1824364460_1, offset: 0, srvID: 0, duration(ms): 419, blockID: 1073741829_1005, type:LAST_IN_PIPELINE
2024-11-18 08:49:09 508 INFO dataNode.DataNode: Receiving BP->1418316472->172.17.0.2:1731919319553:blk_1073741829_1006, duration(ms): 1006, blockID: 1073741829_1006, type:LAST_IN_PIPELINE terminating
2024-11-18 08:49:09 508 INFO hdfs.StateChange: DIR* completeFile: /user/hadoop/input/httpfs-site.xml._COPYING_ is closed by DFSClient_NONMAPREDUCE_-1824364460_1
2024-11-18 08:49:09 530 INFO dataNode.DataNode: Receiving BP->1418316472->172.17.0.2:1731919319553:blk_1073741830_1006, duration(ms): 1006, blockID: 1073741830_1006, type:LAST_IN_PIPELINE terminating
2024-11-18 08:49:09 533 INFO dataNode.DataNode.clienttrace: src: /172.17.0.2:41248 dest: /172.17.0.2:9866 bytes: 3518, op: HDFS_WRITE, clid: DFSClient_NONMAPREDUCE_-1824364460_1, offset: 0, duration(ms): 3810431
2024-11-18 08:49:09 545 INFO dataNode.DataNode: Receiving BP->1418316472->172.17.0.2:1731919319553:blk_1073741830_1006, duration(ms): 3810431
2024-11-18 08:49:09 545 INFO dataNode.DataNode.clienttrace: src: /172.17.0.2:41248 dest: /172.17.0.2:9866 bytes: 3518, op: HDFS_WRITE, clid: DFSClient_NONMAPREDUCE_-1824364460_1, offset: 0, duration(ms): 3810431
2024-11-18 08:49:09 549 INFO hdfs.StateChange: DIR* completeFile: /user/hadoop/input/mes-acls.xml._COPYING_ is closed by DFSClient_NONMAPREDUCE_-1824364460_1
```

```
state = RUNNING [=configuredState]
nodeLocalPriority = 40
nodeLocalityAdditionalDelay = -1
labels=
reservationsContinuedLooking = true
reservationDisabled = true
defaultAppPriorityPerQueue = 0
priority = 0
maxLifetime = -1 seconds
defaultLifetime = -1 seconds
2024-11-18 08:47:43,666 INFO capacity.CapacitySchedulerQueueManager: Initialized queue: default: capacity=1.0, absoluteCapacity=1.0, usedResources=<memory:0, vCores:0>, usedCapacity=0.0, a
bsoluteUsedCapacity=0.0, numContainers=0, effectiveMinResource=<memory:0, vCores:0>, effectiveMaxResource=<memory:0, vCores:0>
2024-11-18 08:47:43,670 INFO capacity.CapacitySchedulerQueueManager: Initialized queue: root: numChildQueue= 1, capacity=1.0, absoluteCapacity=1.0, usedResources=<memory:0, vCores:0>usedCa
pacity=0.0, numApps=0, numContainers=0
2024-11-18 08:47:43,670 INFO capacity.CapacitySchedulerQueueManager: Initialized root queue root: numChildQueue= 1, capacity=1.0, absoluteCapacity=1.0, usedResources=<memory:0, vCores:0>us
edCapacity=0.0, numApps=0, numContainers=0
2024-11-18 08:47:43,671 INFO placement.UserGroupMappingPlacementRule: Initialized queue mappings, override: false
2024-11-18 08:47:43,673 INFO placement.MultiNodeSchedulingManager: MultiNode scheduling is 'false', and configured policies are
2024-11-18 08:47:43,673 INFO capacity.CapacityScheduler: Initialized CapacityScheduler with calculator=<class org.apache.hadoop.yarn.util.resource.DefaultResourceCalculator, minimumAllocati
on=<memory:0, vCores:0>, maximumAllocationMemory: 8192, vCores:4>, asynchronousScheduling=false, asyncScheduleInterval=5ms, multiModePlacementEnabled=false
2024-11-18 08:47:43,676 INFO config.Configuration: dynamic resource allocation is not used
2024-11-18 08:47:43,681 INFO resourcemanager.AMSProcessingChain: Initializing AMS Processing chain. Root Processor=<org.apache.hadoop.yarn.server.resourcemanager.DefaultAMSP
rocessor>.
2024-11-18 08:47:43,681 INFO resourcemanager.ApplicationMasterService: disabled placement handler will be used, all scheduling requests will be rejected.
2024-11-18 08:47:43,685 INFO resourcemanager.AMSProcessingChain: Adding [org.apache.hadoop.yarn.server.resourcemanager.scheduler.constraint.processor.DisabledPlacementProcessor] tp top of
AMS Processing chain.
2024-11-18 08:47:43,696 INFO resourcemanager.ResourceManager: TimelineServicePublisher is not configured
2024-11-18 08:47:43,795 INFO util.log: Logging initialized @844ms
2024-11-18 08:47:43,800 INFO util.log: Using FileSignerSecretProvider, falling back to use random secrets.
2024-11-18 08:47:43,945 INFO http.HttpRequestLog: Http request log for http.requests.resourcemanager is not defined
2024-11-18 08:47:43,965 INFO http.HttpServer2: Added global filter 'safety' (<class org.apache.hadoop.http.HttpServer2$QuotingInputFilter>)
2024-11-18 08:47:43,973 INFO http.HttpServer2: Added filter RMAuthenticationFilter (<class org.apache.hadoop.yarn.server.security.http.RMAuthenticationFilter>) to context cluster
2024-11-18 08:47:43,974 INFO http.HttpServer2: Added filter RMAuthenticationFilter (<class org.apache.hadoop.yarn.server.security.http.RMAuthenticationFilter>) to context logs
2024-11-18 08:47:43,977 INFO http.HttpServer2: Added filter static_user_filter (<class org.apache.hadoop.http.lib.StaticUserWebFilter$StaticUserFilter>) to context cluster
2024-11-18 08:47:43,977 INFO http.HttpServer2: Added filter static_user_filter (<class org.apache.hadoop.http.lib.StaticUserWebFilter$StaticUserFilter>) to context logs
2024-11-18 08:47:43,986 INFO http.HttpServer2: adding path spec: /cluster/*
2024-11-18 08:47:43,986 INFO http.HttpServer2: adding path spec: /ws/*
2024-11-18 08:47:43,987 INFO http.HttpServer2: adding path spec: /app/*
2024-11-18 08:47:44,491 INFO webapp.WebApps: Registered webapp guice modules
2024-11-18 08:47:44,499 INFO http.HttpServer2: Jetty bound to port 8088
2024-11-18 08:47:44,500 INFO server.Server: Jetty-9-3.2.20.v20160605, build timestamp: 2018-06-05T17:11:56Z, git hash: 84205aa28f1la4f3lf2a3b86d1bba2cc8ab69827
2024-11-18 08:47:45,335 INFO server.AuthenticationFilter: Unable to initialize FileSignerSecretProvider, falling back to use random secrets
2024-11-18 08:47:45,336 INFO delegation.AbstractDelegationTokenSecretManager: Updating the current master key for generating delegation tokens
2024-11-18 08:47:45,350 INFO delegation.AbstractDelegationTokenSecretManager: Starting token rotation cycle. The current token will be valid for 60 min(s)
2024-11-18 08:47:45,550 INFO delegation.AbstractDelegationTokenSecretManager: Updating the current master key for generating delegation tokens
2024-11-18 08:47:45,576 INFO handler.ContextHandler: Started o.e.j.s.ServletContextHandler@1caaf3d{/,logs,file:///opt/hadoop-3.2.1/logs/,AVAILABLE}
2024-11-18 08:47:45,579 INFO handler.ContextHandler: Started o.e.j.s.ServletContextHandler@0e01f9b0{/static,jar:file:/opt/hadoop-3.2.1/share/hadoop/yarn/common-3.2.1.jar!/webap
ps/static,AVAILABLE}
```

```
root@80cd224a83dd:/# hadoop jar $HADOOP_HOME/share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar wordcount /user/hadoop/input /user/hadoop/output
```

```

root@889cd22Ha3dd:/# hadoop jar $SHADOOP_HOME/share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar wordcount /user/hadoop/input /user/hadoop/output
2024-11-18 08:49:22.659 INFO impl.MetricsConfig: Scheduled Metrics snapshot at 10 second(s).
2024-11-18 08:49:22.782 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2024-11-18 08:49:23.774 INFO input.FileInputFormat: Total input files to process : 9
2024-11-18 08:49:23.384 INFO mapred.JobSubmitter: number of splits:9
2024-11-18 08:49:23.469 INFO mapred.JobSubmitter: Submitting tokens for job: job_local1972624303_0001
2024-11-18 08:49:23.469 INFO mapred.JobSubmitter: Executing with tokens: []
2024-11-18 08:49:23.585 INFO mapred.Job: The url to track the job: http://localhost:8080/
2024-11-18 08:49:23.585 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2024-11-18 08:49:23.585 INFO mapred.LocalJobRunner: OutputCommitter Algorithm version is 2
2024-11-18 08:49:23.608 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary Folders under output directory:false, ignore cleanup failures: false
2024-11-18 08:49:23.601 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2024-11-18 08:49:23.663 INFO mapred.LocalJobRunner: Waiting for map tasks
2024-11-18 08:49:23.663 INFO mapred.LocalJobRunner: Starting task: attempt_local1972624303_0001_m_000000_0
2024-11-18 08:49:23.713 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2024-11-18 08:49:23.714 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2024-11-18 08:49:23.762 INFO mapred.Task: Using ResourceCalculatorProcessorTree : [ ]
2024-11-18 08:49:23.774 INFO mapred.MapTask: Processing split: hdfs://0.0.0.0:8020/user/hadoop/input/hadoop-policy.xml?r=0+11392
2024-11-18 08:49:23.889 INFO mapred.MapTask: (EQUATOR) 0 kv=26214396(104857584)
2024-11-18 08:49:23.891 INFO mapred.MapTask: mapreduce.task.tracker.address: 0.0.0.0:100
2024-11-18 08:49:23.881 INFO mapred.MapTask: soft limit at 83886080
2024-11-18 08:49:23.881 INFO mapred.MapTask: bufsstart = 0; bufvoid = 104857600
2024-11-18 08:49:23.881 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2024-11-18 08:49:23.891 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2024-11-18 08:49:23.952 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2024-11-18 08:49:24.147 INFO mapred.LocalJobRunner:
2024-11-18 08:49:24.151 INFO mapred.MapTask: Starting flush of map output
2024-11-18 08:49:24.151 INFO mapred.MapTask: Spilling map output
2024-11-18 08:49:24.151 INFO mapred.MapTask: bufsstart = 0; bufvoid = 15783; bufend = 15783; kvstart = 26214396(104857584); kvend = 26209160(104836640); length = 5237/6553600
2024-11-18 08:49:24.151 INFO mapred.MapTask: Map output flushed till 6
2024-11-18 08:49:24.210 INFO mapred.Task: Task:attempt_local1972624303_0001_m_000000_0 is done. And is in the process of committing
2024-11-18 08:49:24.216 INFO mapred.LocalJobRunner: map
2024-11-18 08:49:24.216 INFO mapred.Task: Task 'attempt_local1972624303_0001_m_000000_0' done.
2024-11-18 08:49:24.228 INFO mapred.Task: Final Counters for attempt_local1972624303_0001_m_000000_0: Counters: 24
File System Counters
FILE: Number of bytes read=317864
FILE: Number of bytes written=848370
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=11392
HDFS: Number of bytes written=0
HDFS: Number of read operations=5
HDFS: Number of large read operations=0
HDFS: Number of write operations=1

```

```

2024-11-18 08:49:23.762 INFO mapred.Task: Using ResourceCalculatorProcessTree : []
2024-11-18 08:49:23.771 INFO mapred.MapTask: Processing split: hdfs://88cd224a83dd:8020/user/hadoop/input/hadoop-policy.xml:0+11392
2024-11-18 08:49:23.888 INFO mapred.MapTask: {EXECUTOR} 0 kvi 26214396(104857584)
2024-11-18 08:49:23.881 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2024-11-18 08:49:23.881 INFO mapred.MapTask: soft limit at 83886080
2024-11-18 08:49:23.881 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2024-11-18 08:49:23.881 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2024-11-18 08:49:23.952 INFO mapred.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2024-11-18 08:49:23.952 INFO mapred.LocalJobRunner: Starting flush of map output
2024-11-18 08:49:24.151 INFO mapred.MapTask: Spilling map output
2024-11-18 08:49:24.151 INFO mapred.MapTask: bufstart = 0; bufend = 15783; bufvoid = 104857600
2024-11-18 08:49:24.151 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26209160(104836640); length = 5237/6553600
2024-11-18 08:49:24.151 INFO mapred.MapTask: Finished spill 0
2024-11-18 08:49:24.216 INFO mapred.Task: Task attempt_local1972624303_0001_m_000000_0 is done. And is in the process of committing
2024-11-18 08:49:24.216 INFO mapred.LocalJobRunner: map
2024-11-18 08:49:24.216 INFO mapred.Task: Task 'attempt_local1972624303_0001_m_000000_0' done.
2024-11-18 08:49:24.228 INFO mapred.Task: Final Counters for attempt_local1972624303_0001_m_000000_0: Counters: 24
File System Counters
  FILE: Number of bytes read=317964
  FILE: Number of bytes written=848370
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=11392
  HDFS: Number of bytes written=0
  HDFS: Number of read operations=5
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=1
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=265
  Map output records=1310
  Map output bytes=15783
  Map output materialized bytes=5170
  Input split bytes=125
  Combine input records=1310
  Combine output records=256
  Spilled Records=256
  Failed Maps=0
  Merged Map outputs=0
  GC time elapsed (ms)=0
  Total committed heap usage (bytes)=273678336
File Input Format Counters
  Bytes Read=11392
2024-11-18 08:49:24.228 INFO mapred.LocalJobRunner: Finishing task: attempt_local1972624303_0001_m_000000_0
2024-11-18 08:49:24.238 INFO mapred.LocalJobRunner: Starting task: attempt_local1972624303_0001_m_000001_0
2024-11-18 08:49:24.234 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2024-11-18 08:49:24.234 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false

```

```

WRONG_REDUCE=0
File Output Format Counters
  Bytes Written=10327
2024-11-18 08:49:25.997 INFO mapred.LocalJobRunner: Finishing task: attempt_local1972624303_0001_r_000000_0
2024-11-18 08:49:25.997 INFO mapred.LocalJobRunner: reduce task executor complete.
2024-11-18 08:49:26.045 INFO hdfs.StateChange: DIR* completeFile: /user/hadoop/output/_SUCCESS is closed by DFSClient_NONMAPREDUCE_-1817690479_1
2024-11-18 08:49:26.602 INFO mapreduce.Job: map 100% reduce 100%
2024-11-18 08:49:26.603 INFO mapreduce.Job: Job job_local1972624303_0001 completed successfully
2024-11-18 08:49:26.638 INFO mapreduce.Job: Counters: 36
File System Counters
  FILE: Number of bytes read=3264130
  FILE: Number of bytes written=8618451
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=245235
  HDFS: Number of bytes written=18327
  HDFS: Number of read operations=143
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=12
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=751
  Map output records=3063
  Map output bytes=38519
  Map output materialized bytes=21327
  Input split bytes=1104
  Combine input records=3063
  Combine output records=1223
  Reduce input groups=56
  Reduce output bytes=21327
  Reduce input records=1223
  Reduce output records=568
  Spilled Records=2446
  Shuffled Maps=9
  Failed Shuffles=0
  Merged Map outputs=9
  GC time elapsed (ms)=48
  Total committed heap usage (bytes)=6696206336
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_PARTITION=0
  File Input Format Counters
  Bytes Read=28589
  File Output Format Counters
  Bytes Written=10327

```

t@80cd224a83dd:/# hdfs dfs -cat /user/hadoop/output/part-r-00000

```
root@80cd224a83dd:/# hdfs dfs -cat /user/hadoop/output/part-r-00000
2024-11-18 08:50:18,523 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
"#" 21
"AS" 9
"license"); 9
"alice,bob" 21
"clipping" 1
"(ASF)" 1
(croot 1
(the 9
--> 18
-1 1
-1, 1
0.. 1
1-MAX_INT. 1
1
1
1.. 1
1.. 1
2.. 9
40 2
40+20=60 1
: 2
<!-- 18
</configuration> 9
</description> 33
</name>
</property> 57
<xm 8
<xm-stylesheet 4
<configuration> 9
<description> 31
<description>ACL 25
<description>Default
<name>default.key.acl.DECRYPT_EEK</name> 1
<name>default.key.acl.GENERATE_EEK</name> 1
<name>default.key.acl.MANAGEMENT</name> 1
<name>default.key.acl.READ</name> 1
<name>hadoop.Wms.acl.CREATE</name> 1
<name>hadoop.Wms.acl.DECRYPT_EEK</name> 1
<name>hadoop.Wms.acl.DELETE</name> 1
<name>hadoop.Wms.acl.GENERATE_EEK</name> 1
<name>hadoop.Wms.acl.GET</name> 1
<name>hadoop.Wms.acl.GET_KEYSYNCHRONIZATION</name> 1
<name>hadoop.Wms.acl.ROLLOVER</name> 1
<name>hadoop.Wms.acl.SET_KEY_MATERIAL</name> 1
<name>security.admin.operations.protocol.acl</name> 1
<name>security.applicationclient.protocol.acl</name> 1
<name>security.applicationhistory.protocol.acl</name> 1

<name>security.applicationclient.protocol.acl</name> 1
<name>security.applicationhistory.protocol.acl</name> 1
<name>security.applicationmaster-nodemanager.applicationmaster.protocol.acl</name> 1
<name>security.applicationmaster.protocol.acl</name> 1
<name>security.client.datanode.protocol.acl</name> 1
<name>security.client.protocol.acls</name> 1
<name>security.collector-nodemanager.protocol.acl</name> 1
<name>security.containermanagement.protocol.acl</name> 1
<name>security.datanode.protocol.acl</name> 1
<name>security.distributedscheduling.protocol.acl</name> 1
<name>security.ha.service.protocol.acl</name> 1
<name>security.inter.datanode.protocol.acl</name> 1
<name>security.inter.nodemanager.protocol.acl</name> 1
<name>security.job.allocation.protocol.acl</name> 1
<name>security.job.task.protocol.acl</name> 1
<name>security.job.client.protocol.acl</name> 1
<name>security.namenode.protocol.acl</name> 1
<name>security.qjournal.service.protocol.acl</name> 1
<name>security.refresh.policy.protocol.acl</name> 1
<name>security.refresh.user.mappings.protocol.acl</name> 1
<name>security.resourcelocalizer.protocol.acl</name> 1
<name>security.resourcemanager-administration.protocol.acl</name> 1
<name>security.resourcemanager-configuration.protocol.acl</name> 1
<name>security.zkfc.protocol.acl</name> 1
<name>yarn.scheduler.capacity.application.fail-fast</name> 1
<name>yarn.scheduler.capacity.maximum-am-resource-percent</name> 1
<name>yarn.scheduler.capacity.maximum-applications</name> 1
<name>yarn.scheduler.capacity.node-locality-delays</name> 1
<name>yarn.scheduler.capacity.per-node-heartbeat.maximum-offswitch-assignments</name> 1
<name>yarn.scheduler.capacity.queue-mappings-override.enable</name> 1
<name>yarn.scheduler.capacity.queue-mappings</name> 1
<name>yarn.scheduler.capacity.rack-locality-additional-delay</name> 1
<name>yarn.scheduler.capacity.root.default.acl_administer_queues</name> 1
<name>yarn.scheduler.capacity.root.default.acl_application_max_priority</name> 1
<name>yarn.scheduler.capacity.root.default.acl_submit_applications</name> 1
<name>yarn.scheduler.capacity.root.default.capacity</name> 1
<name>yarn.scheduler.capacity.root.default.default-application-lifetime 1
<name>yarn.scheduler.capacity.root.default.maximum-application-lifetime 1
<name>yarn.scheduler.capacity.root.default.maximum-capacity</name> 1
<name>yarn.scheduler.capacity.root.default.state</name> 1
<name>yarn.scheduler.capacity.root.default.user-limit-factor</name> 1
<name>yarn.scheduler.capacity.root.queues</name> 1
<property>
<property><name>dfs.client.use.datanode.hostname</name><value>true</value></property> 1
<property><name>dfs.datanode.use.datanode.hostname</name><value>true</value></property> 1
<property><name>dfs.datanode.http-bind-host</name><value>0.0.0.0</value></property> 1
<property><name>dfs.datanode.https-bind-host</name><value>0.0.0.0</value></property> 1
<property><name>dfs.datanode.name.dir</name><value>file:///hadoop/dfs/names</value></property> 1
<property><name>dfs.datanode.rpc-bind-host</name><value>0.0.0.0</value></property> 1
```

```
allowed 1
allowed.</description> 21
also 1
an 11
and 70
any 1
applicable 11
application 8
applications 5
applications' 1
applications. 1
approximately 1
are 27
as 10
assign 1
assigning 1
assignments 3
at 10
attempt 1
attempts 2
based 1
basis 1
be 16
blank. 21
block 1
by 52
calculated 1
can 9
can't 1
capacity 1
capacity.</description> 1
case, 1
changes 1
client 2
client-to-datanode 1
clients 3
cluster 3
cluster. 1
code 1
collector 1
comma-separated 21
commands 1
commands. 2
communicate 2
communicate. 12
communicate. 2
compare 2
compliance 9
concurrent 1
config 1
```

```
submit 2
submitted 4
submit. 1
syntax 1
taken 2
taken. 1
target 1
tasks 1
tasktracker. 1
than 3
that 7
the 139
then 2
this 28
time 1
timeline 3
timestamp. 1
to 53
too 1
type="text/xsl" 4
u:User:User 1
under 28
unique 1
updating 1
use 11
used 24
user 48
user. 2
users? 1
users. 24
users, wheel". 21
uses 2
using 1
v2 1
valid. 1
value 30
values 2
values? 2
version=1.0" 5
version=1.0"?> 3
via 1
well 1
when 3
which 10
while 1
who 3
will 12
with 28
work 1
writing, 9
```

```
root@80cd224a83dd:/# exit
```

```
exit
```

```
C:\Users\pirai>docker cp 80cd224a83dd:/opt/hadoop-3.2.1/etc/hadoop/core-site.xml ./core-site.xml
```

```
C:\Users\pirai>docker cp 80cd224a83dd:/opt/hadoop-3.2.1/etc/hadoop/hdfs-site.xml ./hdfs-site.xml
```

```
root@80cd224a83dd:/# 2024-11-18 08:57:14,505 INFO localizer.ResourceLocalizationService: Cache Size Before Clean: 0, Total Deleted: 0, Public Deleted: 0, Private Deleted: 0
2024-11-18 08:57:46,286 INFO scheduler.AbstractYarnScheduler: Release request cache is cleaned up
2024-11-18 09:07:14,445 INFO localizer.ResourceLocalizationService: Cache Size Before Clean: 0, Total Deleted: 0, Public Deleted: 0, Private Deleted: 0
root@80cd224a83dd:/# docker-compose up -d
bash: docker-compose: command not found
root@80cd224a83dd:/# 2024-11-18 09:17:14,392 INFO localizer.ResourceLocalizationService: Cache Size Before Clean: 0, Total Deleted: 0, Public Deleted: 0, Private Deleted: 0
root@80cd224a83dd:/# nano $HADOOP_HOME/etc/hadoop/core-site.xml
bash: nano: command not found
root@80cd224a83dd:/# nano $HADOOP_HOME/etc/hadoop/hdfs-site.xml
bash: nano: command not found
root@80cd224a83dd:/# 2024-11-18 09:27:14,359 INFO localizer.ResourceLocalizationService: Cache Size Before Clean: 0, Total Deleted: 0, Public Deleted: 0, Private Deleted: 0
root@80cd224a83dd:/# docker cp 80cd224a83dd:/opt/hadoop-3.2.1/etc/hadoop/core-site.xml ./core-site.xml
bash: docker: command not found
root@80cd224a83dd:/# 2024-11-18 09:37:14,310 INFO localizer.ResourceLocalizationService: Cache Size Before Clean: 0, Total Deleted: 0, Public Deleted: 0, Private Deleted: 0
root@80cd224a83dd:/# exit
exit
C:\Users\pirai>docker cp 80cd224a83dd:/opt/hadoop-3.2.1/etc/hadoop/core-site.xml ./core-site.xml
Successfully copied 2.56kB to C:\Users\pirai\core-site.xml
C:\Users\pirai>docker cp 80cd224a83dd:/opt/hadoop-3.2.1/etc/hadoop/hdfs-site.xml ./hdfs-site.xml
Successfully copied 3.07kB to C:\Users\pirai\hdfs-site.xml
C:\Users\pirai>
```